
SOFTWARE REQUIREMENTS SPECIFICATION

for

CS 4ZP6 Capstone Project

Version 0.0

Prepared by Brendan Duke, Andrew Kohnen,
Udip Patel, David Pitkanen, Jordan Viveiros

Using Volere Template Edition 13

McMaster Text to Motion Database

October 12, 2016

Contents

1	Project Drivers	6
1.1	The Purpose of the Project	6
1.1.1	The User Business or Background of the Project Effort	6
1.1.2	Goals of the Project	6
1.2	The Client, the Customer, and Other Stakeholders	7
1.2.1	The Client	7
1.2.2	The Customer	7
1.2.3	Other Stakeholders	7
1.3	Users of the Product	8
1.3.1	The Hands-on Users of the Product	8
1.3.2	Priorities Assigned to Users	8
1.3.3	User Participation	8
1.3.4	Maintenance Users and Service Technicians	8
2	Project Constraints	10
2.1	Mandated Constraints	10
2.1.1	Solution Constraints	10
2.1.2	Implementation Environment of the Current System	11
2.1.3	Partner or Collaborative Applications	12
2.1.4	Off-the-Shelf Software	12
2.1.5	Anticipated Workplace Environment	13
2.1.6	Schedule Constraints	13
2.1.7	Budget Constraints	14
2.2	Naming Conventions and Definitions	14
2.2.1	Definitions of All Terms, Including Acronyms, Used in the Project	14
2.2.2	Data Dictionary for any Included Models	15
2.3	Relevant Facts and Assumptions	15
2.3.1	Facts	15
2.3.2	Assumptions	16
3	Functional Requirements	17
3.1	The Scope of the Work	17
3.1.1	The Current Situation	17
3.1.2	The Context of the Work	17
3.1.3	Work Partitioning	17
3.2	The Scope of the Product	20
3.2.1	Product Boundary	20

3.2.2	Product Use-case List	20
3.2.3	Individual Product Use Cases	20
3.3	Functional and Data Requirements	20
3.3.1	Functional Requirements	20
3.3.2	Data Requirements	20
4	Nonfunctional Requirements	21
4.1	Look and Feel Requirements	23
4.1.1	Appearance Requirements	23
4.1.2	Style Requirements	23
4.2	Usability and Humanity Requirements	23
4.2.1	Ease of Use Requirements	23
4.2.2	Personalization and Internationalization Requirements	23
4.2.3	Learning Requirements	23
4.2.4	Understandability and Politeness Requirements	23
4.2.5	Accessibility Requirements	23
4.3	Performance Requirements	23
4.3.1	Speed and Latency Requirements	23
4.3.2	Safety-Critical Requirements	23
4.3.3	Precision or Accuracy Requirements	23
4.3.4	Reliability and Availability Requirements	23
4.3.5	Robustness or Fault-Tolerance Requirements	23
4.3.6	Capacity Requirements	23
4.3.7	Scaling of Extensibility Requirements	23
4.3.8	Longevity Requirements	23
4.4	Operational and Environmental Requirements	23
4.4.1	Expected Physical Environment	23
4.4.2	Requirements for Interfacing with Adjacent Systems	23
4.4.3	Productization Requirements	23
4.4.4	Release Requirements	23
4.5	Maintainability and Support Requirements	23
4.5.1	Maintenance Requirements	23
4.5.2	Supportability Requirements	23
4.5.3	Adaptability Requirements	23
4.6	Security Requirements	23
4.6.1	Access Requirements	23
4.6.2	Integrity Requirements	23
4.6.3	Privacy Requirements	23
4.6.4	Audit Requirements	23
4.6.5	Immunity Requirements	23
4.7	Cultural and Political Requirements	23
4.7.1	Cultural Requirements	23
4.7.2	Political Requirements	23

4.8	Legal Requirements	23
4.8.1	Compliance Requirements	23
4.8.2	Standards Requirements	23
5	Project Issues	24
5.1	Open Issues	26
5.2	Off-the-Shelf Solutions	26
5.2.1	Ready-Made Products	26
5.2.2	Reusable Components	26
5.2.3	Products That Can Be Copied	26
5.3	New Problems	26
5.3.1	Effects on the Current Environment	26
5.3.2	Effects on the Installed Systems	26
5.3.3	Potential User Problems	26
5.3.4	Limitations in the Anticipated Implementation Environment That May Inhibit the New Product	26
5.3.5	Follow-Up Problems	26
5.4	Tasks	26
5.4.1	Project Planning	26
5.4.2	Planning of the Development Phases	26
5.5	Migration to the New Product	26
5.5.1	Requirements for Migration of the New Product	26
5.5.2	Data That Has to Be Modified or Translated for the New System .	26
5.6	Risks	26
5.7	Costs	26
5.8	User Documentation and Training	26
5.8.1	User Documentation Requirements	26
5.8.2	Training Requirements	26
5.9	Waiting Room	26
5.10	Ideas for Solutions	26

Revision History

Name	Date	Reason For Changes	Version
Brendan Duke	Oct. 7th, 2016	Initial Version	0.0

1 Project Drivers

1.1 The Purpose of the Project

1.1.1 The User Business or Background of the Project Effort

There is an existing project at the University of Guelph that aims to create a system for “computational storytelling”. The goal of the computational storytelling project is to create a system that takes as input a basic story composed of five sentences, and outputs an animated movie based on the story, which is produced in collaboration between an AI and human director.

As an initial step in the computational storytelling project, the University of Guelph group requires a database of “human motion” that is stored with rich text annotations. Such a database is required as a source of training data for the computational storytelling project to use in their methods to convert text to animated motion.

No satisfactory database of human motion data that is stored with associated text descriptions exists currently. However, there are existing databases of videos of people doing various actions with accompanying text describing those actions, for example the Charades database or MSR-VTT.

Our McMaster group has been approached to assist in this initial step in the computational storytelling project in two ways. Firstly, we are to develop software, based on existing research, that is able to process video and derive human motion data (e.g. joint positions over time) from the video. Secondly, we are to utilize data from an existing database that already has text annotations, such as Charades, using our video-to-motion processing software to generate a new database that contains both rich text annotations and motion data.

1.1.2 Goals of the Project

The goal of this project is to create a database, web-interface to said database, and a deployable software bundle providing access to already-established human pose estimation methods. Creating this database, website and software suite will allow the larger text-to-motion project to use the relationships between motion data and text annotations developed through the pose estimation software in order to provide a pose and word pairing, which can be used for animation.

1.2 The Client, the Customer, and Other Stakeholders

1.2.1 The Client

The current clients for this project are Dr. Taylor and his graduate student Thor Jonsson. Dr. Taylor is the primary driver to develop a website and database where annotated motion information can be generated and pulled from as a growth point into the larger text to motion project. They will be using the database to train Recurrent Neural Networks (RNNs) that will pair actions and their pose found within the database to words or combinations found in the input story.

1.2.2 The Customer

The customers are included within the clients since building this database and website combination will be utilized by Dr. Taylor's research team and their external partners. In addition to Dr. Taylor and his research team this project would appeal to anyone that needed a pairing of actions and pose estimations as the website would be readily available to others.

In general, customers of the product will be researchers in the machine learning community who are interested in multi-modal learning, and specifically in systems that link text to human motion. Said customers will have a high degree of knowledge related to machine learning theory. However, they cannot be assumed to have a high degree of skill in any programming language with a steep skill curve, such as C++ or Haskell. Also, the customer is unlikely to be willing to invest a large amount of time in learning how to use the software produced by the McMaster Text-to-Motion project.

1.2.3 Other Stakeholders

Other stakeholders affected by the project include Dr. He, our group's internal supervisor and teacher of the CS 4ZP6 Capstone Project course, and the team members of our group.

Dr. He is a professor at McMaster who may not have the same specialized research knowledge as members of Dr. Taylor's group. Dr. He requires an explanation of all aspects of the project, as she will be responsible for assigning a grade to the entire group. Dr. He will require updates on the progress of the group in the form of deliverables that are part of the CS 4ZP6 syllabus.

The members of our CS 4ZP6 capstone group, namely Brendan Duke, Andrew Kohnen, Udip Patel, David Pitkanen and Jordan Viveiros, are also stakeholders affected by the project. For the most part, our group members did not have any specialized knowledge related to the project before commencing, although that knowledge is being acquired as the project progresses. The group members will require full involvement in all aspects of the project, as well as supporting knowledge and direction from Thor Jonsson and Dr. Taylor.

1.3 Users of the Product

1.3.1 The Hands-on Users of the Product

User Category	User Role	Subject Matter Experience	Technological Experience
Dr. Taylor's Group	Using the database to train an RNN to create animations from text.	Master	Master
Other machine learning researchers	Using the product for any multi-modal machine learning use-case involving text and human motion.	Master	Journeyman. This user category cannot be assumed to have a high degree of skill in complex programming languages such as C++.
Amateur machine learning enthusiasts	Using the McMaster Text-to-Motion database and software suite to learn about multi-modal machine learning and human pose estimation.	Journeyman	Journeyman

1.3.2 Priorities Assigned to Users

Our **key users** are members of Dr. Taylor's research group. **Secondary users** are other members of the machine learning community. Amateur machine learning enthusiasts are **unimportant users**.

1.3.3 User Participation

Thor Jonsson and Dr. Taylor will be expected to assist in supporting our group with their domain knowledge of deep learning methods. They will also be expected to participate in shaping the interfaces to the product (both the web interface and the programming interface to the database) by using the prototypes of those interfaces and providing feedback.

The minimum amount of participation from Dr. Taylor and Thor would be participation in a meeting with our group members on a bi-weekly to monthly basis, as well as participating in weekly correspondence electronically (e.g. by e-mail).

1.3.4 Maintenance Users and Service Technicians

Maintenance users would certainly be members of Dr. Taylor's research group, as they will be using the software produced by the project after its completion and may need to add changes to the product.

Once the product is open-sourced into the community, maintenance users could range from machine learning researchers to amateur machine learning enthusiasts. These users could be expected to fix bugs or add new features that were not in the initial scope of the project.

2 Project Constraints

Note that constraints also use the requirement shell, and therefore also use the requirement numbering.

Customer satisfaction and dissatisfaction have been evaluated in the "Priority" entry.

Also there are no conflicts between requirements, or supporting materials, so these entries have been dropped.

2.1 Mandated Constraints

2.1.1 Solution Constraints

Req. #: 1	Req. Type: 4a. Solution Constraint	Use Case #: 3
Description: The human pose estimation component should use deep learning methods.		
Rationale: This constraint is to allow Dr. Taylor's group to integrate the software into their existing text-to-motion pipeline.		
Originator: Dr. Graham Taylor		
Fit Criterion: Dr. Taylor should confirm that the deep learning methods used in the human pose estimator are satisfactory.		
Priority: High		History: Created September 26th, 2016.

Req. #: 2	Req. Type: 4a. Solution Constraint	Use Case #: 4
Description: Use a standard format such as LMDB or HDF5 for storing text-motion data.		
Rationale: Having the data in a standard format will enable users to re-use existing code to manipulate that data.		
Originator: Thor Jonsson		
Fit Criterion: Run a set of existing tests to manipulate the standard data format (e.g. LMDB) and assert that those tests must pass.		
Priority: High	History: Created October 3rd, 2016.	

2.1.2 Implementation Environment of the Current System

Req. #: 3	Req. Type: 4b. Implementation Environment	Use Case #: 3
Description: The Text-to-Motion Software Suite must run under Linux.		
Rationale: Linux is the operating system used by the Guelph Machine Learning research lab, and also the most commonly used operating system in the research community.		
Originator: Dr. Graham Taylor		
Fit Criterion: Automated builds and testing should pass on popular Linux distributions: Ubuntu, Fedora and RHEL.		
Priority: High	History: Created September 26th, 2016.	

Req. #: 4	Req. Type: 4b. Implementation Environment	Use Case #: 3
Description: Major APIs to the Text-to-Motion database must be accessible from the Python programming language.		
Rationale: Python is the language used by the rest of Dr. Taylor's text-to-motion pipeline. Python is a popular, easy-to-use, and quick-to-prototype language, and is therefore one of the most favoured programming languages among the Machine Learning research community.		
Originator: Dr. Graham Taylor		
Fit Criterion: There must be hooks to all major interfaces written in Python, and there must be tests that are directly testing the Python interfaces.		
Priority: High		History: Created September 26th, 2016.

2.1.3 Partner or Collaborative Applications

The McMaster Text-to-Motion database software will collaborate with the rest of the University of Guelph group's text-to-motion pipeline. While their application is still in a research stage, we can expect that their application will be written in Python.

Furthermore, the animation component of Dr. Taylor's group's project will be using the proprietary software application Muvizu.

2.1.4 Off-the-Shelf Software

There are a number of off-the-shelf libraries that can and should be used by this project in order to implement the requirements. In particular, deep learning libraries should be used for the training of neural networks and for doing numerical computations.

Two notable deep learning frameworks are listed below.

Caffe, originally developed by the Berkeley Vision and Learning Center (BVLC), is a deep learning framework written in C++ with Python and Matlab wrappers.

Caffe has a strong ConvNet implementation and is popular amongst the computer vision community. However, Caffe's RNN implementation and support for language models is lacking compared with other libraries.

Caffe is released under the BSD 2-Clause license.

TensorFlow, originally developed by researchers at Google, is a deep learning and numerical computation framework written in C++ and Python.

TensorFlow has a clean, modular architecture and since it provides both Python and C++ interfaces, code can be prototyped in a rich high-level interpreted language before

deployment in a high-performance, scalable environment.

TensorFlow is open sourced under the Apache 2.0 open source license.

It will also be necessary to use a library for accessing video codecs. This is in order to convert video streams to and from sets of images, which will be input to the human pose estimation software.

FFmpeg is a leading multimedia framework that is able to decode and encode video, in addition to having many other features. FFmpeg can be utilized for any video encoding and decoding purposes in the project.

To implement search-by-text functionality into the McMaster Text-to-Motion Database, it will be helpful to leverage an existing software package that implements full-text search.

Sphinx is such a software package, as it implements full-text search. Sphinx also integrates well with SQL databases.

Sphinx is licensed under GPL version 2.

2.1.5 Anticipated Workplace Environment

As the McMaster Text-to-Motion Database is a software product, there are no requirements considerations specifically pertaining to the anticipated workplace environment of users of the product. Any considerations of environment, such as operating system environment, have already been covered under the “Implementation Environment” section.

2.1.6 Schedule Constraints

Req. #: 5	Req. Type: 4f. Schedule Constraint	Use Case #: 3
Description: A demonstration of the product must be completed by February 13th, 2017.		
Rationale: This is a deadline that is part of the CS 4ZP6 course.		
Originator: Dr. Wen Bo He		
Fit Criterion: By February 13th, 2017, fully functional versions of the web interface, database interface and human pose estimation software must all be working as verified by Dr. He.		
Priority: High	History: Created September 26th, 2016.	

Req. #: 6	Req. Type: 4f. Schedule Constraint	Use Case #: 3
Description: The project must be completed by April 5th, 2017.		
Rationale: The project is part of the CS 4ZP6 Capstone Project course.		
Originator: Dr. Wen Bo He		
Fit Criterion: All documentation, testing and implementation must be completed and checked in to GitHub by April 5th, 2017.		
Priority: High		History: Created September 21th, 2016.

2.1.7 Budget Constraints

There is no budget allocated for this project. As such, hardware requirements such as GPUs to carry out the computations required for the project must be provided by the client, Dr. Taylor, or borrowed through other means such as seeking help from McMaster professors.

2.2 Naming Conventions and Definitions

2.2.1 Definitions of All Terms, Including Acronyms, Used in the Project

The Project when used, is referring to the McMaster Text to Motion Database project. The project aims to generate a database of human pose estimation model information that is linked to videos of human motion containing rich text annotations.

Human Pose Estimation is the process of estimating the configuration, or pose, of the body based on a single still image or a sequence of images that comprise a video. Human pose estimation may find the chin, radius, humerus, and other bone and joint positions.

Charades is a dataset composed of approximately 10K videos of daily indoor activities, complete with associated action-describing sentences, collected through Amazon Mechanical Turk[1].

MSR-VTT, standing for “Microsoft Research Video to Text”, is a large-scale video benchmark for the task of translating video to text. MSR-VTT provides 10K video clips spanning 41.2 hours and containing 200K clip-sentence pairs in total[2].

Feedforward Neural Networks are artificial neural networks where connections between the units do *not* form a cycle). They are the simplest type of neural network, because information moves in only one direction.

ConvNets or **Convolutional Neural Networks** are a type of feed-forward artificial neural network. ConvNets are inspired by the visual cortex and are commonly used in

visual recognition applications.

RNNs or **Recurrent Neural Networks** are a class of artificial neural networks where units form a directed cycle, in contrast with feed-forward neural networks.

Deep Belief Networks are a type of deep neural network composed of multiple layers of “hidden units” (variables that are not observable), with connections between layers but not between units of a given layer.

Multi-modal neural language models are models of natural language that can be conditioned on other modalities, e.g. high-level image features[3].

Muvizu is an interactive 3D animation package designed for quick storytelling.

An **autoencoder** is an artificial neural network used to learn a representation of a set of data. The simplest form of an autoencoder is a feed-forward neural net with an input layer, an output layer, and one or more hidden layers. Instead of being trained to predict an output Y from input X , an autoencoder is trained to reconstruct its own input X' from X .

2.2.2 Data Dictionary for any Included Models

Here we define the different information flows displayed in Figure 3.1, the work context diagram.

A **skeleton overlay** refers to a set of human joint and body-part estimation data that has been associated with a given video stream. A skeleton overlay may also encompass an actual graphical stick-figure, displayed on top of the video stream, with circles indicating the position of each joint.

Text annotated video refers to a video that has a time-stamped natural-language description associated with it. Similarly, **text annotated motion data** refers to human pose data that is time-stamped and coupled with a natural-language description of the actions comprised by those human pose data.

2.3 Relevant Facts and Assumptions

2.3.1 Facts

- The training of deep neural networks can take on the order of days to weeks. For example, training two of the different neural network architectures used in [4] on four NVIDIA GTX Titan GPUs took 3 and 7 days, respectively.
- ConvNets take a fixed size input and generate a fixed size output, whereas RNNs can handle arbitrary input/output lengths. This is one reason why ConvNets are more popular in vision applications, while RNNs are used for language applications. However, neural network architectures can be used that combine both ConvNets and RNNs.
- The uncompressed size of the Charades dataset is 55GB at full resolution, and 16GB at 480p.

- The goals of this project, specifically generating a database of text-annotated human motion data, constitute new research and there are no existing solutions available.

2.3.2 Assumptions

- It has been assumed that it is possible to re-use existing software components in order to achieve the human pose estimation use-case, as it would take a full eight months just to implement an algorithm such as that of [4] from scratch.
- It has been assumed that this project will not encompass generating a dataset of videos with text annotations, and that we will be able to rely on an existing dataset such as the Charades dataset.
- It is assumed that a storage format for human pose estimation data already exists, and will be available for this project to leverage.

3 Functional Requirements

3.1 The Scope of the Work

3.1.1 The Current Situation

There is a large amount of existing research into human pose estimation, which this project will leverage. Based on Constraint 1, we focus on existing solutions that use deep learning methods.

[4] present a ConvNet architecture for human pose estimation from videos, which is able to benefit from temporal context across multiple frames using optical flow. This work is focused on upper-body human pose estimation only.

[5] propose a ConvNet model for predicting 2D human body poses in an image. This model is able to achieve state-of-the-art results using a simple architecture, and draws on the work done in [4].

[6] introduces *Convolutional Pose Machines (CPMs)* for pose estimation in images. CPMs consist of a sequence of ConvNets that iteratively produce 2D belief maps. Source code from this paper is available, and makes use of the Caffe deep learning library.

In [7], the authors estimate 3D full-body human poses from a monocular image sequence. They train a ConvNet to predict the locations of joints in 2D, then estimate human poses in 3D via an Expectation-Maximization algorithm over the entire video sequence.

On the animation side of the larger computational storytelling problem our project is assisting in solving, [8] use deep learning to synthesize character movements based on high-level parameters such as a curve that the character should follow. The learned motion is represented by the hidden units of a convolutional autoencoder.

3.1.2 The Context of the Work

The context diagram, displaying adjacent systems to this project, is shown in Figure 3.1.

3.1.3 Work Partitioning

The events/use-cases for the project are laid out in Table 3.1.

Figure 3.1: McMaster Text-to-Motion Database Work Context Diagram

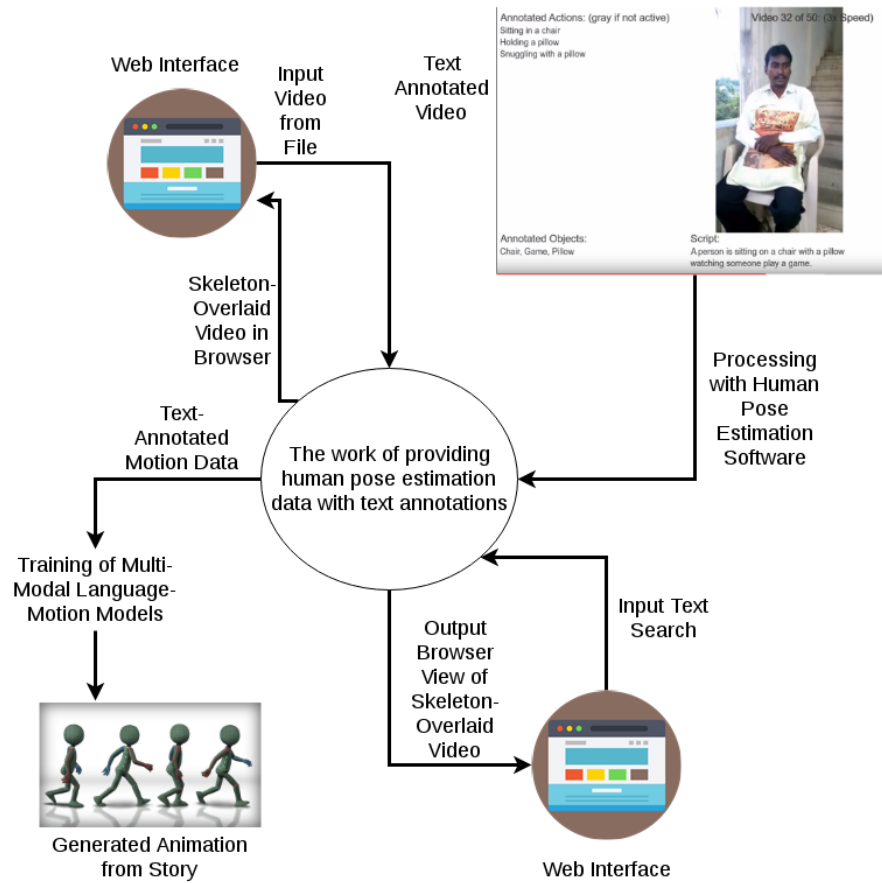


Table 3.1: Business Event List

Event #	Event Name	Input and Output	Summary
1.	Web Interface Skeleton Overlay	IN: An image or video with humans in it. OUT: The same image or video, with a skeleton overlaid on top of all humans indicating their bone and joint positions.	Allow users to observe the human pose estimation component in real time through a web interface.
2.	Web Interface Text-to-Motion	IN: Word or phrase describing a human pose or action. OUT: Rich-text-annotated video corresponding to the input word/phrase, complete with overlaid skeleton.	Allow users to see the output of searches on the database using pose and/or action keywords, such as “run” or “kneeling”.
3.	Software Interface Skeleton Overlay	IN: A stream of video with humans depicted. OUT: A set of human pose estimations corresponding to the video, in a standard data format.	Users should be able to use the human pose estimation solution to generate their own motion data set.
4.	Database Interface Text-to-Motion	IN: Word or phrase describing a human pose or action. OUT: Video in common encoding (e.g. MP4), associated rich-text-annotations, and human pose estimations in a standardized format.	Provide users direct access to the raw motion-estimation data format based on action-keyword database lookup.

3.2 The Scope of the Product

3.2.1 Product Boundary

We do not include a product boundary diagram, and instead refer to Figure 3.1.

3.2.2 Product Use-case List

3.2.3 Individual Product Use Cases

3.3 Functional and Data Requirements

3.3.1 Functional Requirements

Req. #: 7	Req. Type: 9a. Functional Requirement	Use Case #: 3
Description: The text-to-motion software suite will provide an API to read individual frames in RGB format from a video stream. At least MP4, MP2 and AAC must be supported.		
Rationale: Researchers may wish to do their own processing on RGB frames before feeding those frames into the human pose estimation module.		
Originator: For a given set of test video streams, the frame-capture API must produce RGB frames identical to known reference frames.		
Fit Criterion: Moderate		
Priority: Created October 5th, 2016.		History: Brendan Duke

3.3.2 Data Requirements

4 Nonfunctional Requirements

4.1 Look and Feel Requirements

4.1.1 Appearance Requirements

4.1.2 Style Requirements

4.2 Usability and Humanity Requirements

4.2.1 Ease of Use Requirements

4.2.2 Personalization and Internationalization Requirements

4.2.3 Learning Requirements

4.2.4 Understandability and Politeness Requirements

4.2.5 Accessibility Requirements

4.3 Performance Requirements

4.3.1 Speed and Latency Requirements

4.3.2 Safety-Critical Requirements

4.3.3 Precision or Accuracy Requirements

4.3.4 Reliability and Availability Requirements

4.3.5 Robustness or Fault-Tolerance Requirements

4.3.6 Capacity Requirements

4.3.7 Scaling of Extensibility Requirements

4.3.8 Longevity Requirements

4.4 Operational and Environmental Requirements

4.4.1 Expected Physical Environment

4.4.2 Requirements for Interfacing with Adjacent Systems

4.4.3 Productization Requirements

4.4.4 Release Requirements

4.5 Maintainability and Support Requirements

4.5.1 Maintenance Requirements

4.5.2 Supportability Requirements

4.5.3 Adaptability Requirements

23

4.6 Security Requirements

4.6.1 Access Requirements

4.6.2 Integrity Requirements

4.6.3 Privacy Requirements

4.6.4 Audit Requirements

5 Project Issues

5.1 Open Issues

5.2 Off-the-Shelf Solutions

5.2.1 Ready-Made Products

5.2.2 Reusable Components

5.2.3 Products That Can Be Copied

5.3 New Problems

5.3.1 Effects on the Current Environment

5.3.2 Effects on the Installed Systems

5.3.3 Potential User Problems

5.3.4 Limitations in the Anticipated Implementation Environment That May Inhibit the New Product

5.3.5 Follow-Up Problems

5.4 Tasks

5.4.1 Project Planning

5.4.2 Planning of the Development Phases

5.5 Migration to the New Product

5.5.1 Requirements for Migration of the New Product

5.5.2 Data That Has to Be Modified or Translated for the New System

5.6 Risks

5.7 Costs

5.8 User Documentation and Training

5.8.1 User Documentation Requirements

5.8.2 Training Requirements

5.9 Waiting Room

5.10 Ideas for Solutions

Bibliography

- [1] Charades dataset. Allen Institute for Artificial Intelligence. [Online]. Available: <http://allenai.org/plato/charades/>
- [2] J. Xu, T. Mei, T. Yao, and Y. Rui, “Msr-vtt: A large video description dataset for bridging video and language.” IEEE International Conference on Computer Vision and Pattern Recognition (CVPR), June 2016. [Online]. Available: <https://www.microsoft.com/en-us/research/publication/msr-vtt-a-large-video-description-dataset-for-bridging-video-and-language/>
- [3] R. Kiros, R. Salakhutdinov, and R. S. Zemel, “Unifying visual-semantic embeddings with multimodal neural language models,” *CoRR*, vol. abs/1411.2539, 2014. [Online]. Available: <http://arxiv.org/abs/1411.2539>
- [4] T. Pfister, J. Charles, and A. Zisserman, “Flowing convnets for human pose estimation in videos,” *CoRR*, vol. abs/1506.02897, 2015. [Online]. Available: <http://arxiv.org/abs/1506.02897>
- [5] V. Belagiannis and A. Zisserman, “Recurrent human pose estimation,” *CoRR*, vol. abs/1605.02914, 2016. [Online]. Available: <http://arxiv.org/abs/1605.02914>
- [6] S. Wei, V. Ramakrishna, T. Kanade, and Y. Sheikh, “Convolutional pose machines,” *CoRR*, vol. abs/1602.00134, 2016. [Online]. Available: <http://arxiv.org/abs/1602.00134>
- [7] X. Zhou, M. Zhu, K. Derpanis, and K. Daniilidis, “Sparseness meets deepness: 3d human pose estimation from monocular video,” in *CVPR*, 2016.
- [8] D. Holden, J. Saito, and T. Komura, “A deep learning framework for character motion synthesis and editing,” *ACM Trans. Graph.*, vol. 35, no. 4, pp. 138:1–138:11, Jul. 2016. [Online]. Available: <http://doi.acm.org/10.1145/2897824.2925975>