

---

# SOFTWARE REQUIREMENTS SPECIFICATION

for

CS 4ZP6 Capstone Project

Version 0.0

Prepared by Brendan Duke, Andrew Kohnen,  
Udip Patel, David Pitkanen, Jordan Viveiros

Using Volere Template Edition 13

McMaster Text to Motion Database

October 12, 2016

# Contents

<b>1</b>	<b>Project Drivers</b>	<b>6</b>
1.1	The Purpose of the Project . . . . .	6
1.1.1	The User Business or Background of the Project Effort . . . . .	6
1.1.2	Goals of the Project . . . . .	6
1.2	The Client, the Customer, and Other Stakeholders . . . . .	7
1.2.1	The Client . . . . .	7
1.2.2	The Customer . . . . .	7
1.2.3	Other Stakeholders . . . . .	7
1.3	Users of the Product . . . . .	8
1.3.1	The Hands-on Users of the Product . . . . .	8
1.3.2	Priorities Assigned to Users . . . . .	8
1.3.3	User Participation . . . . .	8
1.3.4	Maintenance Users and Service Technicians . . . . .	8
<b>2</b>	<b>Project Constraints</b>	<b>10</b>
2.1	Mandated Constraints . . . . .	10
2.1.1	Solution Constraints . . . . .	10
2.1.2	Implementation Environment of the Current System . . . . .	11
2.1.3	Partner or Collaborative Applications . . . . .	12
2.1.4	Off-the-Shelf Software . . . . .	12
2.1.5	Anticipated Workplace Environment . . . . .	13
2.1.6	Schedule Constraints . . . . .	13
2.1.7	Budget Constraints . . . . .	14
2.2	Naming Conventions and Definitions . . . . .	14
2.2.1	Definitions of All Terms, Including Acronyms, Used in the Project . . . . .	14
2.2.2	Data Dictionary for any Included Models . . . . .	15
2.3	Relevant Facts and Assumptions . . . . .	15
2.3.1	Facts . . . . .	15
2.3.2	Assumptions . . . . .	16
<b>3</b>	<b>Functional Requirements</b>	<b>17</b>
3.1	The Scope of the Work . . . . .	17
3.1.1	The Current Situation . . . . .	17
3.1.2	The Context of the Work . . . . .	17
3.1.3	Work Partitioning . . . . .	17
3.2	The Scope of the Product . . . . .	20
3.2.1	Product Boundary . . . . .	20

3.2.2	Product Use-case List . . . . .	20
3.2.3	Individual Product Use Cases . . . . .	20
3.3	Functional and Data Requirements . . . . .	21
3.3.1	Functional Requirements . . . . .	21
3.3.2	Data Requirements . . . . .	23
<b>4</b>	<b>Nonfunctional Requirements</b>	<b>24</b>
4.1	Look and Feel Requirements . . . . .	24
4.1.1	Appearance Requirements . . . . .	24
4.1.2	Style Requirements . . . . .	25
4.2	Usability and Humanity Requirements . . . . .	25
4.2.1	Ease of Use Requirements . . . . .	25
4.2.2	Personalization and Internationalization Requirements . . . . .	26
4.2.3	Learning Requirements . . . . .	27
4.2.4	Understandability and Politeness Requirements . . . . .	28
4.2.5	Accessibility Requirements . . . . .	29
4.3	Performance Requirements . . . . .	30
4.3.1	Speed and Latency Requirements . . . . .	30
4.3.2	Safety-Critical Requirements . . . . .	32
4.3.3	Precision or Accuracy Requirements . . . . .	32
4.3.4	Reliability and Availability Requirements . . . . .	32
4.3.5	Robustness or Fault-Tolerance Requirements . . . . .	32
4.3.6	Capacity Requirements . . . . .	32
4.3.7	Scaling of Extensibility Requirements . . . . .	32
4.3.8	Longevity Requirements . . . . .	32
4.4	Operational and Environmental Requirements . . . . .	32
4.4.1	Expected Physical Environment . . . . .	32
4.4.2	Requirements for Interfacing with Adjacent Systems . . . . .	32
4.4.3	Productization Requirements . . . . .	32
4.4.4	Release Requirements . . . . .	32
4.5	Maintainability and Support Requirements . . . . .	32
4.5.1	Maintenance Requirements . . . . .	32
4.5.2	Supportability Requirements . . . . .	32
4.5.3	Adaptability Requirements . . . . .	32
4.6	Security Requirements . . . . .	32
4.6.1	Access Requirements . . . . .	32
4.6.2	Integrity Requirements . . . . .	32
4.6.3	Privacy Requirements . . . . .	32
4.6.4	Audit Requirements . . . . .	32
4.6.5	Immunity Requirements . . . . .	32
4.7	Cultural and Political Requirements . . . . .	32
4.7.1	Cultural Requirements . . . . .	32
4.7.2	Political Requirements . . . . .	32

4.8	Legal Requirements . . . . .	32
4.8.1	Compliance Requirements . . . . .	32
4.8.2	Standards Requirements . . . . .	32
<b>5</b>	<b>Project Issues</b>	<b>33</b>
5.1	Open Issues . . . . .	35
5.2	Off-the-Shelf Solutions . . . . .	35
5.2.1	Ready-Made Products . . . . .	35
5.2.2	Reusable Components . . . . .	35
5.2.3	Products That Can Be Copied . . . . .	35
5.3	New Problems . . . . .	35
5.3.1	Effects on the Current Environment . . . . .	35
5.3.2	Effects on the Installed Systems . . . . .	35
5.3.3	Potential User Problems . . . . .	35
5.3.4	Limitations in the Anticipated Implementation Environment That May Inhibit the New Product . . . . .	35
5.3.5	Follow-Up Problems . . . . .	35
5.4	Tasks . . . . .	35
5.4.1	Project Planning . . . . .	35
5.4.2	Planning of the Development Phases . . . . .	35
5.5	Migration to the New Product . . . . .	35
5.5.1	Requirements for Migration of the New Product . . . . .	35
5.5.2	Data That Has to Be Modified or Translated for the New System .	35
5.6	Risks . . . . .	35
5.7	Costs . . . . .	35
5.8	User Documentation and Training . . . . .	35
5.8.1	User Documentation Requirements . . . . .	35
5.8.2	Training Requirements . . . . .	35
5.9	Waiting Room . . . . .	35
5.10	Ideas for Solutions . . . . .	35

## Revision History

Name	Date	Reason For Changes	Version
Brendan Duke	Oct. 7th, 2016	Initial Version	0.0

# 1 Project Drivers

## 1.1 The Purpose of the Project

### 1.1.1 The User Business or Background of the Project Effort

There is an existing project at the University of Guelph that aims to create a system for “computational storytelling”. The goal of the computational storytelling project is to create a system that takes as input a basic story composed of five sentences, and outputs an animated movie based on the story, which is produced in collaboration between an AI and human director.

As an initial step in the computational storytelling project, the University of Guelph group requires a database of “human motion” that is stored with rich text annotations. Such a database is required as a source of training data for the computational storytelling project to use in their methods to convert text to animated motion.

No satisfactory database of human motion data that is stored with associated text descriptions exists currently. However, there are existing databases of videos of people doing various actions with accompanying text describing those actions, for example the Charades database or MSR-VTT.

Our McMaster group has been approached to assist in this initial step in the computational storytelling project in two ways. Firstly, we are to develop software, based on existing research, that is able to process video and derive human motion data (e.g. joint positions over time) from the video. Secondly, we are to utilize data from an existing database that already has text annotations, such as Charades, using our video-to-motion processing software to generate a new database that contains both rich text annotations and motion data.

### 1.1.2 Goals of the Project

The goal of this project is to create a database, web-interface to said database, and a deployable software bundle providing access to already-established human pose estimation methods. Creating this database, website and software suite will allow the larger text-to-motion project to use the relationships between motion data and text annotations developed through the pose estimation software in order to provide a pose and word pairing, which can be used for animation.

## **1.2 The Client, the Customer, and Other Stakeholders**

### **1.2.1 The Client**

The current clients for this project are Dr. Taylor and his graduate student Thor Jonsson. Dr. Taylor is the primary driver to develop a website and database where annotated motion information can be generated and pulled from as a growth point into the larger text to motion project. They will be using the database to train Recurrent Neural Networks (RNNs) that will pair actions and their pose found within the database to words or combinations found in the input story.

### **1.2.2 The Customer**

The customers are included within the clients since building this database and website combination will be utilized by Dr. Taylor's research team and their external partners. In addition to Dr. Taylor and his research team this project would appeal to anyone that needed a pairing of actions and pose estimations as the website would be readily available to others.

In general, customers of the product will be researchers in the machine learning community who are interested in multi-modal learning, and specifically in systems that link text to human motion. Said customers will have a high degree of knowledge related to machine learning theory. However, they cannot be assumed to have a high degree of skill in any programming language with a steep skill curve, such as C++ or Haskell. Also, the customer is unlikely to be willing to invest a large amount of time in learning how to use the software produced by the McMaster Text-to-Motion project.

### **1.2.3 Other Stakeholders**

Other stakeholders affected by the project include Dr. He, our group's internal supervisor and teacher of the CS 4ZP6 Capstone Project course, and the team members of our group.

Dr. He is a professor at McMaster who may not have the same specialized research knowledge as members of Dr. Taylor's group. Dr. He requires an explanation of all aspects of the project, as she will be responsible for assigning a grade to the entire group. Dr. He will require updates on the progress of the group in the form of deliverables that are part of the CS 4ZP6 syllabus.

The members of our CS 4ZP6 capstone group, namely Brendan Duke, Andrew Kohnen, Udip Patel, David Pitkanen and Jordan Viveiros, are also stakeholders affected by the project. For the most part, our group members did not have any specialized knowledge related to the project before commencing, although that knowledge is being acquired as the project progresses. The group members will require full involvement in all aspects of the project, as well as supporting knowledge and direction from Thor Jonsson and Dr. Taylor.

## 1.3 Users of the Product

### 1.3.1 The Hands-on Users of the Product

User Category	User Role	Subject Matter Experience	Technological Experience
Dr. Taylor's Group	Using the database to train an RNN to create animations from text.	Master	Master
Other machine learning researchers	Using the product for any multi-modal machine learning use-case involving text and human motion.	Master	Journeyman. This user category cannot be assumed to have a high degree of skill in complex programming languages such as C++.
Amateur machine learning enthusiasts	Using the McMaster Text-to-Motion database and software suite to learn about multi-modal machine learning and human pose estimation.	Journeyman	Journeyman

### 1.3.2 Priorities Assigned to Users

Our **key users** are members of Dr. Taylor's research group. **Secondary users** are other members of the machine learning community. Amateur machine learning enthusiasts are **unimportant users**.

### 1.3.3 User Participation

Thor Jonsson and Dr. Taylor will be expected to assist in supporting our group with their domain knowledge of deep learning methods. They will also be expected to participate in shaping the interfaces to the product (both the web interface and the programming interface to the database) by using the prototypes of those interfaces and providing feedback.

The minimum amount of participation from Dr. Taylor and Thor would be participation in a meeting with our group members on a bi-weekly to monthly basis, as well as participating in weekly correspondence electronically (e.g. by e-mail).

### 1.3.4 Maintenance Users and Service Technicians

Maintenance users would certainly be members of Dr. Taylor's research group, as they will be using the software produced by the project after its completion and may need to add changes to the product.



Once the product is open-sourced into the community, maintenance users could range from machine learning researchers to amateur machine learning enthusiasts. These users could be expected to fix bugs or add new features that were not in the initial scope of the project.

## 2 Project Constraints

Note that constraints also use the requirement shell, and therefore also use the requirement numbering.

Customer satisfaction and dissatisfaction have been evaluated in the "Priority" entry.

Also there are no conflicts between requirements, or supporting materials, so these entries have been dropped.

### 2.1 Mandated Constraints

#### 2.1.1 Solution Constraints

<b>Req. #:</b> 1	<b>Req. Type:</b> 2.1.1 Solution Constraint	<b>Use Case #:</b> 3
<b>Description:</b> The human pose estimation component should use deep learning methods.		
<b>Rationale:</b> This constraint is to allow Dr. Taylor's group to integrate the software into their existing text-to-motion pipeline.		
<b>Originator:</b> Dr. Graham Taylor		
<b>Fit Criterion:</b> Dr. Taylor should confirm that the deep learning methods used in the human pose estimator are satisfactory.		
<b>Priority:</b> High		<b>History:</b> Created September 26th, 2016.

<b>Req. #:</b> 2	<b>Req. Type:</b> 2.1.1 Solution Constraint	<b>Use Case #:</b> 4
<b>Description:</b> Use a standard format such as LMDB or HDF5 for storing text-motion data.		
<b>Rationale:</b> Having the data in a standard format will enable users to re-use existing code to manipulate that data.		
<b>Originator:</b> Thor Jonsson		
<b>Fit Criterion:</b> Run a set of existing tests to manipulate the standard data format (e.g. LMDB) and assert that those tests must pass.		
<b>Priority:</b> High	<b>History:</b> Created October 3rd, 2016.	

### 2.1.2 Implementation Environment of the Current System

<b>Req. #:</b> 3	<b>Req. Type:</b> 2.1.2 Implementation Environment	<b>Use Case #:</b> 3
<b>Description:</b> The Text-to-Motion Software Suite must run under Linux.		
<b>Rationale:</b> Linux is the operating system used by the Guelph Machine Learning research lab, and also the most commonly used operating system in the research community.		
<b>Originator:</b> Dr. Graham Taylor		
<b>Fit Criterion:</b> Automated builds and testing should pass on popular Linux distributions: Ubuntu, Fedora and RHEL.		
<b>Priority:</b> High	<b>History:</b> Created September 26th, 2016.	

<b>Req. #:</b> 4	<b>Req. Type:</b> 2.1.2 Implementation Environment	<b>Use Case #:</b> 3
<b>Description:</b> Major APIs to the Text-to-Motion database must be accessible from the Python programming language.		
<b>Rationale:</b> Python is the language used by the rest of Dr. Taylor's text-to-motion pipeline. Python is a popular, easy-to-use, and quick-to-prototype language, and is therefore one of the most favoured programming languages among the Machine Learning research community.		
<b>Originator:</b> Dr. Graham Taylor		
<b>Fit Criterion:</b> There must be hooks to all major interfaces written in Python, and there must be tests that are directly testing the Python interfaces.		
<b>Priority:</b> High		<b>History:</b> Created September 26th, 2016.

### 2.1.3 Partner or Collaborative Applications

The McMaster Text-to-Motion database software will collaborate with the rest of the University of Guelph group's text-to-motion pipeline. While their application is still in a research stage, we can expect that their application will be written in Python.

Furthermore, the animation component of Dr. Taylor's group's project will be using the proprietary software application Muvizu.

### 2.1.4 Off-the-Shelf Software

There are a number of off-the-shelf libraries that can and should be used by this project in order to implement the requirements. In particular, deep learning libraries should be used for the training of neural networks and for doing numerical computations.

Two notable deep learning frameworks are listed below.

**Caffe**, originally developed by the Berkeley Vision and Learning Center (BVLC), is a deep learning framework written in C++ with Python and Matlab wrappers.

Caffe has a strong ConvNet implementation and is popular amongst the computer vision community. However, Caffe's RNN implementation and support for language models is lacking compared with other libraries.

Caffe is released under the BSD 2-Clause license.

**TensorFlow**, originally developed by researchers at Google, is a deep learning and numerical computation framework written in C++ and Python.

TensorFlow has a clean, modular architecture and since it provides both Python and C++ interfaces, code can be prototyped in a rich high-level interpreted language before

deployment in a high-performance, scalable environment.

TensorFlow is open sourced under the Apache 2.0 open source license.

It will also be necessary to use a library for accessing video codecs. This is in order to convert video streams to and from sets of images, which will be input to the human pose estimation software.

**FFmpeg** is a leading multimedia framework that is able to decode and encode video, in addition to having many other features. FFmpeg can be utilized for any video encoding and decoding purposes in the project.

To implement search-by-text functionality into the McMaster Text-to-Motion Database, it will be helpful to leverage an existing software package that implements full-text search.

**Sphinx** is such a software package, as it implements full-text search. Sphinx also integrates well with SQL databases.

Sphinx is licensed under GPL version 2.

### 2.1.5 Anticipated Workplace Environment

As the McMaster Text-to-Motion Database is a software product, there are no requirements considerations specifically pertaining to the anticipated workplace environment of users of the product. Any considerations of environment, such as operating system environment, have already been covered under the “Implementation Environment” section.

### 2.1.6 Schedule Constraints

<b>Req. #:</b> 5	<b>Req. Type:</b> 2.1.6 Schedule Constraint	<b>Use Case #:</b> 3
<b>Description:</b> A demonstration of the product must be completed by February 13th, 2017.		
<b>Rationale:</b> This is a deadline that is part of the CS 4ZP6 course.		
<b>Originator:</b> Dr. Wen Bo He		
<b>Fit Criterion:</b> By February 13th, 2017, fully functional versions of the web interface, database interface and human pose estimation software must all be working as verified by Dr. He.		
<b>Priority:</b> High	<b>History:</b> Created September 26th, 2016.	

<b>Req. #:</b> 6	<b>Req. Type:</b> 2.1.6 Schedule Constraint	<b>Use Case #:</b> 3
<b>Description:</b> The project must be completed by April 5th, 2017.		
<b>Rationale:</b> The project is part of the CS 4ZP6 Capstone Project course.		
<b>Originator:</b> Dr. Wen Bo He		
<b>Fit Criterion:</b> All documentation, testing and implementation must be completed and checked in to GitHub by April 5th, 2017.		
<b>Priority:</b> High		<b>History:</b> Created September 21th, 2016.

### 2.1.7 Budget Constraints

There is no budget allocated for this project. As such, hardware requirements such as GPUs to carry out the computations required for the project must be provided by the client, Dr. Taylor, or borrowed through other means such as seeking help from McMaster professors.

## 2.2 Naming Conventions and Definitions

### 2.2.1 Definitions of All Terms, Including Acronyms, Used in the Project

**The Project** when used, is referring to the McMaster Text to Motion Database project. The project aims to generate a database of human pose estimation model information that is linked to videos of human motion containing rich text annotations.

**Human Pose Estimation** is the process of estimating the configuration, or pose, of the body based on a single still image or a sequence of images that comprise a video. Human pose estimation may find the chin, radius, humerus, and other bone and joint positions.

**Charades** is a dataset composed of approximately 10K videos of daily indoor activities, complete with associated action-describing sentences, collected through Amazon Mechanical Turk[1].

**MSR-VTT**, standing for “Microsoft Research Video to Text”, is a large-scale video benchmark for the task of translating video to text. MSR-VTT provides 10K video clips spanning 41.2 hours and containing 200K clip-sentence pairs in total[2].

**Feedforward Neural Networks** are artificial neural networks where connections between the units do *not* form a cycle). They are the simplest type of neural network, because information moves in only one direction.

**ConvNets** or **Convolutional Neural Networks** are a type of feed-forward artificial neural network. ConvNets are inspired by the visual cortex and are commonly used in

visual recognition applications.

**RNNs** or **Recurrent Neural Networks** are a class of artificial neural networks where units form a directed cycle, in contrast with feed-forward neural networks.

**Deep Belief Networks** are a type of deep neural network composed of multiple layers of “hidden units” (variables that are not observable), with connections between layers but not between units of a given layer.

**Multi-modal neural language models** are models of natural language that can be conditioned on other modalities, e.g. high-level image features[3].

**Muvizu** is an interactive 3D animation package designed for quick storytelling.

An **autoencoder** is an artificial neural network used to learn a representation of a set of data. The simplest form of an autoencoder is a feed-forward neural net with an input layer, an output layer, and one or more hidden layers. Instead of being trained to predict an output  $Y$  from input  $X$ , an autoencoder is trained to reconstruct its own input  $X'$  from  $X$ .

**Ground truth** refers to the expected classification, i.e. the expected output in a machine learning classification problem, based on a given input.

**CMU Graphics Lab Motion Capture Database** is a dataset of human motions that were collected at the Motion Capture Lab of Carnegie Mellon University. The motion capture data was collected by recording the movements of humans wearing a body suit with 41 markers taped on. The 3D motion data in this dataset is available in two formats (.vsk/.v or .asf/.amc), which describe a *skeleton and its joints* and *movement data*, respectively.

## 2.2.2 Data Dictionary for any Included Models

Here we define the different information flows displayed in Figure 3.1, the work context diagram.

A **skeleton overlay** refers to a set of human joint and body-part estimation data that has been associated with a given video stream. A skeleton overlay may also encompass an actual graphical stick-figure, displayed on top of the video stream, with circles indicating the position of each joint.

**Text annotated video** refers to a video that has a time-stamped natural-language description associated with it. Similarly, **text annotated motion data** refers to human pose data that is time-stamped and coupled with a natural-language description of the actions comprised by those human pose data.

## 2.3 Relevant Facts and Assumptions

### 2.3.1 Facts

- The training of deep neural networks can take on the order of days to weeks. For example, training two of the different neural network architectures used in [4] on four NVIDIA GTX Titan GPUs took 3 and 7 days, respectively.

- ConvNets take a fixed size input and generate a fixed size output, whereas RNNs can handle arbitrary input/output lengths. This is one reason why ConvNets are more popular in vision applications, while RNNs are used for language applications. However, neural network architectures can be used that combine both ConvNets and RNNs.
- The uncompressed size of the Charades dataset is 55GB at full resolution, and 16GB at 480p.
- The goals of this project, specifically generating a database of text-annotated human motion data, constitute new research and there are no existing solutions available.

### 2.3.2 Assumptions

- It has been assumed that it is possible to re-use existing software components in order to achieve the human pose estimation use-case, as it would take a full eight months just to implement an algorithm such as that of [4] from scratch.
- It has been assumed that this project will not encompass generating a dataset of videos with text annotations, and that we will be able to rely on an existing dataset such as the Charades dataset.
- It is assumed that a storage format for human pose estimation data already exists, and will be available for this project to leverage.



## 3 Functional Requirements

### 3.1 The Scope of the Work

#### 3.1.1 The Current Situation

There is a large amount of existing research into human pose estimation, which this project will leverage. Based on Constraint 1, we focus on existing solutions that use deep learning methods.

[4] present a ConvNet architecture for human pose estimation from videos, which is able to benefit from temporal context across multiple frames using optical flow. This work is focused on upper-body human pose estimation only.

[5] propose a ConvNet model for predicting 2D human body poses in an image. This model is able to achieve state-of-the-art results using a simple architecture, and draws on the work done in [4].

[6] introduces *Convolutional Pose Machines (CPMs)* for pose estimation in images. CPMs consist of a sequence of ConvNets that iteratively produce 2D belief maps. Source code from this paper is available, and makes use of the Caffe deep learning library.

In [7], the authors estimate 3D full-body human poses from a monocular image sequence. They train a ConvNet to predict the locations of joints in 2D, then estimate human poses in 3D via an Expectation-Maximization algorithm over the entire video sequence.

On the animation side of the larger computational storytelling problem our project is assisting in solving, [8] use deep learning to synthesize character movements based on high-level parameters such as a curve that the character should follow. The learned motion is represented by the hidden units of a convolutional autoencoder.

#### 3.1.2 The Context of the Work

The context diagram, displaying adjacent systems to this project, is shown in Figure 3.1.

#### 3.1.3 Work Partitioning

The events/use-cases for the project are laid out in Table 3.1.

Figure 3.1: McMaster Text-to-Motion Database Work Context Diagram

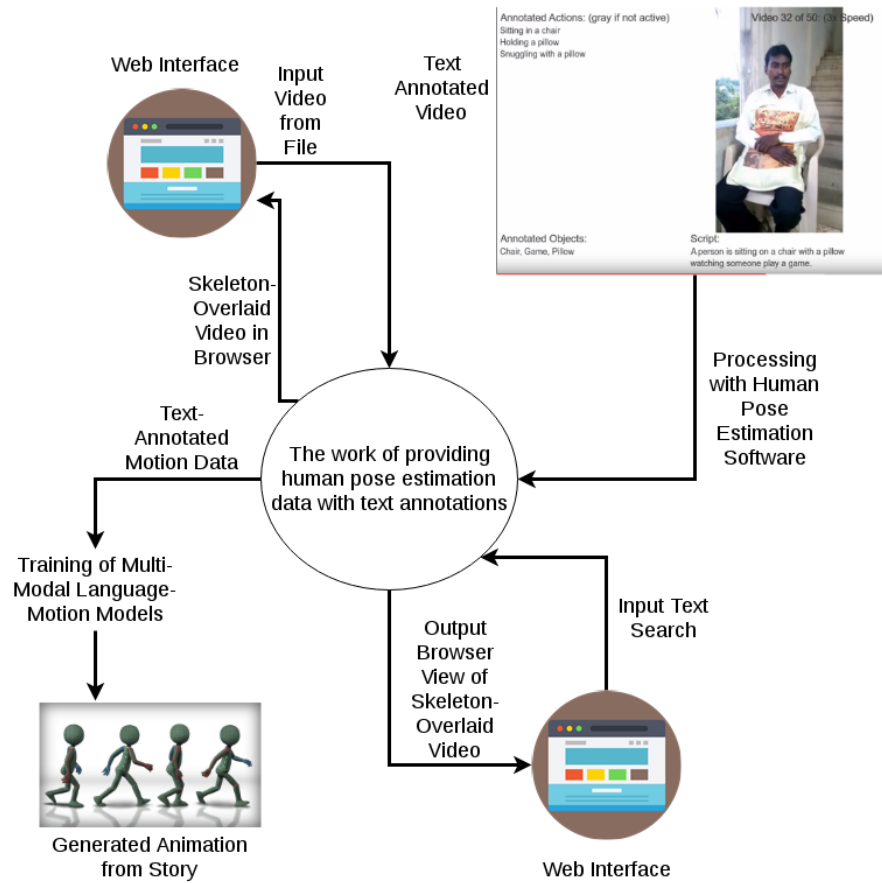


Table 3.1: Business Event List

Event #	Event Name	Input and Output	Summary
1.	Web Interface Skeleton Overlay	<b>IN:</b> An image or video with humans in it. <b>OUT:</b> The same image or video, with a skeleton overlaid on top of all humans indicating their bone and joint positions.	Allow users to observe the human pose estimation component in real time through a web interface.
2.	Web Interface Text-to-Motion	<b>IN:</b> Word or phrase describing a human pose or action. <b>OUT:</b> Rich-text-annotated video corresponding to the input word/phrase, complete with overlaid skeleton.	Allow users to see the output of searches on the database using pose and/or action keywords, such as “run” or “kneeling”.
3.	Software Interface Skeleton Overlay	<b>IN:</b> A stream of video with humans depicted. <b>OUT:</b> A set of human pose estimations corresponding to the video, in a standard data format.	Users should be able to use the human pose estimation solution to generate their own motion data set.
4.	Database Interface Text-to-Motion	<b>IN:</b> Word or phrase describing a human pose or action. <b>OUT:</b> Video in common encoding (e.g. MP4), associated rich-text-annotations, and human pose estimations in a standardized format.	Provide users direct access to the raw motion-estimation data format based on action-keyword database lookup.

## 3.2 The Scope of the Product

### 3.2.1 Product Boundary

We do not include a product boundary diagram, and instead refer to Figure 3.1.

All interfaces to the McMaster Text-to-Motion Database will be software interfaces except the web interfaces, which will have a human component. There will be a search box requiring user input to do text-to-motion lookups, and there will be a graphical web interface for users to upload videos to use the skeleton-overlay functionality.

### 3.2.2 Product Use-case List

The product use cases correspond exactly to the business events list, and therefore to see product use-cases one should refer to Table 3.1.

### 3.2.3 Individual Product Use Cases

Use cases 1 and 2 would be used in a scenario where the user is looking to do experiments with immediate feedback, in order to gain intuition about the product's functionality. E.g. a researcher may want to use the Web Interface Skeleton Overlay feature (use case 1) to visually gauge the human pose estimation software's accuracy before spending the time to process a large amount of video data.

Use cases 3 and 4 could be used in a more large scale experiment. E.g. human motion data could be attached to a large amount of video data using use case 3, or a large amount of text-annotated motion data could be retrieved from the database programmatically using use case 4.

## 3.3 Functional and Data Requirements

### 3.3.1 Functional Requirements

<b>Req. #:</b> 7	<b>Req. Type:</b> 3.3.1 Functional Requirement	<b>Use Case #:</b> 3
<b>Description:</b> The text-to-motion software suite will provide an API to read individual frames in RGB format from a video stream. At least MP4, MP2 and AAC must be supported.		
<b>Rationale:</b> Researchers may wish to do their own processing on RGB frames before feeding those frames into the human pose estimation module.		
<b>Originator:</b> Brendan Duke		
<b>Fit Criterion:</b> For a given set of test video streams, the frame-capture API must produce RGB frames identical to known reference frames.		
<b>Priority:</b> Moderate		<b>History:</b> Created October 5th, 2016.
<b>Req. #:</b> 8	<b>Req. Type:</b> 3.3.1 Functional Requirement	<b>Use Case #:</b> 3
<b>Description:</b> The skeleton overlay software must output a set of human pose data that is time-stamped and accurate to the ground truth joint positions in the frames corresponding to those time stamps.		
<b>Rationale:</b> If the software does not have accurate output, then the output cannot be used to generate human motion.		
<b>Originator:</b> Brendan Duke		
<b>Fit Criterion:</b> A set of double-blind tests should be designed that allows test subjects to compare the human pose data output from the McMaster Text-to-Motion software suite against human motion skeletons rendered from the CMU Graphics Lab Motion Capture Database. The fit criterion would be that the human motion data produced by the McMaster software is statistically indistinguishable from CMU's MoCap data from sensors on actual humans.		
<b>Priority:</b> High		<b>History:</b> Created September 26th, 2016.

<b>Req. #:</b> 9	<b>Req. Type:</b> 3.3.1 Functional Requirement	<b>Use Case #:</b> 4
<p><b>Description:</b> The domain of inputs to the text-to-motion database search should output a range that covers the entire set of videos in the database.</p> <p><b>Rationale:</b> There is text associated with each video, therefore for every video there should be some set of matching text that can retrieve it.</p> <p><b>Originator:</b> Brendan Duke</p> <p><b>Fit Criterion:</b> A test application could be designed to intentionally pull text data from the description of each video and assert that the video from which that text data was obtained appears in the output of the search.</p> <p><b>Priority:</b> High</p> <p><b>History:</b> Created October 10th, 2016.</p>		
<b>Req. #:</b> 10	<b>Req. Type:</b> 3.3.1 Functional Requirement	<b>Use Case #:</b> 4
<p><b>Description:</b> The videos retrieved from a text-to-motion search should not contain any results with text descriptions not containing the search terms.</p> <p><b>Rationale:</b> False positives would make the text-to-motion retrieval difficult to use, since returned videos would not reliably match the search terms.</p> <p><b>Originator:</b> Brendan Duke</p> <p><b>Fit Criterion:</b> A test application could be designed to, for each video, search a random set of terms not in that video's text description and assert that the given video is not returned in the search results.</p> <p><b>Priority:</b> High</p> <p><b>History:</b> Created October 10th, 2016.</p>		

<b>Req. #:</b> 11	<b>Req. Type:</b> 3.3.1 Functional Requirement	<b>Use Case #:</b> 4
<b>Description:</b> The videos retrieved from a text-to-motion search should be ordered by relevance.		
<b>Rationale:</b> Users will want to use the motion data that most closely fits the text they have searched for.		
<b>Originator:</b> Brendan Duke		
<b>Fit Criterion:</b> A smaller set of test data can be created by hand and sorted by a human, and the results can be compared against an expected output, also generated by hand. The output from a larger set of test data could be compared against the output from a known full-text search engine, such as Sphinx, or if we are using Sphinx, then Lucene.		
<b>Priority:</b> High		<b>History:</b> Created October 10th, 2016.

### 3.3.2 Data Requirements

The McMaster Text-to-Motion project is not strictly adhering to an object-oriented design, and therefore there is no class model that must be followed as part of the requirements.

Interactions between classes, or modules, will be diagrammed as part of the module hierarchy during the design iterations of the project development.

## 4 Nonfunctional Requirements

### 4.1 Look and Feel Requirements

#### 4.1.1 Appearance Requirements

<b>Req. #:</b> 12	<b>Req. Type:</b> <a href="#">4.1.1</a> Appearance Requirement	<b>Use Case #:</b> <a href="#">1</a>
<b>Description:</b> The web interface shall have a color-scheme that is focused on contrast.		
<b>Rationale:</b> A color shceme with bright and dark colors clashing is bound to captivate users of the website.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> At least 6 out of every 10 users should be left with a positive impression of the color scheme.		
<b>Priority:</b> Low		<b>History:</b> Created October 8th, 2016.



### 4.1.2 Style Requirements

<b>Req. #:</b> 13	<b>Req. Type:</b> 4.1.2 Style Requirement	<b>Use Case #:</b> 1
<b>Description:</b> The web interface shall be designed to look minimalistic (in terms of text) and visually informative through graphs and other visual aides.		
<b>Rationale:</b> To be able to convey the complex processes of a deep learning model in a simple way that can be understood by any end user, regardless of their level of knowledge. The amount of text displayed should also be minimized to avoid overwhelming a first-time user with too much information.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> The web interface shall be able to display a graph that tracks how the deep learning model analyzes an input image. This is similar to an open-source website called <a href="#">TensorFlow Playground</a> .		
<b>Priority:</b> Low	<b>History:</b> Created October 8th, 2016.	

## 4.2 Usability and Humanity Requirements

### 4.2.1 Ease of Use Requirements

<b>Req. #:</b> 14	<b>Req. Type:</b> 4.2.1 Ease of Use Requirement	<b>Use Case #:</b> 1
<b>Description:</b> The functionalites of the web interface shall be easy to use for any casual internet user above the age of 12 who understands how to upload and download files to a website.		
<b>Rationale:</b> Having a simpler interface can open up this web app to a lot of potential users.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> Any end user at least 12 years of age shall be able know how to upload an image to the web interface within 15 seconds of visiting the website.		
<b>Priority:</b> Low	<b>History:</b> Created October 10th, 2016.	

<b>Req. #:</b> 15	<b>Req. Type:</b> 4.2.1 Ease of Use Requirement	<b>Use Case #:</b> 1
<b>Description:</b> The user shall not be expected to remember anything to be able to interact with the Web Interface.		
<b>Rationale:</b> This is just to reduce the amount of work that the user has to do in order to be able to get the desired result.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> The Web Interface shall only ask for images, videos or text from the user, and without prompting the user again, it will return the given image or video (with or without text annotations).		
<b>Priority:</b> Low	<b>History:</b> Created October 10th, 2016.	

---

<b>Req. #:</b> 16	<b>Req. Type:</b> 4.2.1 Ease of Use Requirement	<b>Use Case #:</b> 2
<b>Description:</b> The user shall be able to type in a descriptive word or phrase into a text box that will be easy to find from the web interface.		
<b>Rationale:</b> This way, the functionality provided by the web application will be made obvious and accessible to the user.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> The end user shall be able to input one word or phrase into the text box (written in HTML/CSS) using their keyboard, and be able to click on a 'Submit' button to view the functionality of the web app.		
<b>Priority:</b> Medium	<b>History:</b> Created October 10th, 2016.	

#### 4.2.2 Personalization and Internationalization Requirements

As of Revision 0, there are no requirements for Language support. Since the web Interface and the Software Interface would rely on Natural Language Processing for English, adding language support for multiple languages would be too large of a feature to implement for now.

### 4.2.3 Learning Requirements

??

<b>Req. #:</b> 17	<b>Req. Type:</b> ?? Learning Requirement	<b>Use Case #:</b> 1
<b>Description:</b> Any end-user shall be able to use the web application and generate a skeleton-overlay without having any prior training.		
<b>Rationale:</b> This is because all the user needs to be able to do is upload and download a file to/from the website.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> Any user that visits the website shall be able to understand how to use the skeleton-overlay functionality within 1 minute of visiting the website.		
<b>Priority:</b> High	<b>History:</b> Created October 10th, 2016.	

---

<b>Req. #:</b> 18	<b>Req. Type:</b> ?? Learning Requirement	<b>Use Case #:</b> 2
<b>Description:</b> Any end-user shall be able to use the web application to instantiate a 'text-to-motion' event without having any prior training.		
<b>Rationale:</b> This is because all the user needs to do is type in a descriptive word or phrase into a text box.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> Any user that visits the website shall be able to understand how to use the 'text-to-motion' functionality within 30 seconds of visiting the website.		
<b>Priority:</b> High	<b>History:</b> Created October 10th, 2016.	

<b>Req. #:</b> 19	<b>Req. Type:</b> ?? Learning Requirement	<b>Use Case #:</b> 3
<b>Description:</b> The functionalities of the software interface shall be easily understood by programmers who have experience with APIs.		
<b>Rationale:</b> This is because using a command line API still takes some base knowledge of software to set up, and we cannot assume that every user will be able to set up our API without any problems even if they follow all of the detailed steps.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> Any user should be able to set up the Software Interface and be able to use its functionalities (described in events 3 and 4) through the command line interface within 15 minutes.		
<b>Priority:</b> Medium		<b>History:</b> Created October 10th, 2016.

#### 4.2.4 Understandability and Politeness Requirements

<b>Req. #:</b> 20	<b>Req. Type:</b> 4.2.4 Understandability and Politeness Requirement	<b>Use Case #:</b> 1
<b>Description:</b> The end user of the web interface shall not be able to see the inner workings of the deep learning model that is used to estimate the human pose or action from an image/video		
<b>Rationale:</b> This is because the user should be obscured from having to deal with the complexities of deep learning. Furthermore, it would be computationally expensive to show the users such low-level detail in a way that they could understand it.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> The end user of the web interface shall be able to get a valid output from the skeleton-overlay function without seeing any of the deep learning code or processes.		
<b>Priority:</b> High		<b>History:</b> Created October 10th, 2016.

<b>Req. #:</b> 21	<b>Req. Type:</b> 4.2.4 Understandability and Politeness Requirement	<b>Use Case #:</b> 4
<b>Description:</b> The end user of the software interface shall not know how the program is able to map a word or phrase to a human pose or action		
<b>Rationale:</b> Using the software interface should be simple, and if we show these highly complex data mappings to the user, we might risk overwhelming the user by giving them a large amount of information that is of very little use.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> The end user of the web interface shall be able to get a valid output from the software interface without being exposed to the database mappings that are used to determine the valid output.		
<b>Priority:</b> High		<b>History:</b> Created October 11th, 2016.

#### 4.2.5 Accessibility Requirements

As of Revision 9, there are no accessibility requirements for this application. Since both interfaces only contain a virtual/visual component and no physical mechanisms, there is no need to provide any accessibility requirements for now

## 4.3 Performance Requirements

### 4.3.1 Speed and Latency Requirements

<b>Req. #:</b> 22	<b>Req. Type:</b> 4.3.1 Speed and Latency Requirement	<b>Use Case #:</b> 2
<b>Description:</b> The web interface shall be able to connect to an external database with no problems and be able to store items and query for them quickly		
<b>Rationale:</b> The speed is necessary here because generating the actual output itself is going to be computationally intensive so the program should not be spend a lot of time waiting on responses from the database.		
<b>Originator:</b> Udip Patel		
<b>Fit Criterion:</b> The web interface should be able to establish an initial connection to the database within a span of 5 seconds. The web interface should be able to transfer a file of size 500MB to the database in around 50 seconds, assuming an upload speed of 10Mbps. The web interface should be able to transfer a file of size 500MB from the database to the user in around 25 seconds, assuming a download speed of 20Mbps.		
<b>Priority:</b> High		<b>History:</b> Created October 11th, 2016.



#### **4.3.2 Safety-Critical Requirements**

#### **4.3.3 Precision or Accuracy Requirements**

#### **4.3.4 Reliability and Availability Requirements**

#### **4.3.5 Robustness or Fault-Tolerance Requirements**

#### **4.3.6 Capacity Requirements**

#### **4.3.7 Scaling of Extensibility Requirements**

#### **4.3.8 Longevity Requirements**

### **4.4 Operational and Environmental Requirements**

#### **4.4.1 Expected Physical Environment**

#### **4.4.2 Requirements for Interfacing with Adjacent Systems**

#### **4.4.3 Productization Requirements**

#### **4.4.4 Release Requirements**

### **4.5 Maintainability and Support Requirements**

#### **4.5.1 Maintenance Requirements**

#### **4.5.2 Supportability Requirements**

#### **4.5.3 Adaptability Requirements**

### **4.6 Security Requirements**

#### **4.6.1 Access Requirements**

#### **4.6.2 Integrity Requirements**

#### **4.6.3 Privacy Requirements**

#### **4.6.4 Audit Requirements**

#### **4.6.5 Immunity Requirements**

### **4.7 Cultural and Political Requirements**

#### **4.7.1 Cultural Requirements**

#### **4.7.2 Political Requirements**

### **4.8 Legal Requirements**

#### **4.8.1 Compliance Requirements**

#### **4.8.2 Standards Requirements**







# **5 Project Issues**

## **5.1 Open Issues**

## **5.2 Off-the-Shelf Solutions**

### **5.2.1 Ready-Made Products**

### **5.2.2 Reusable Components**

### **5.2.3 Products That Can Be Copied**

## **5.3 New Problems**

### **5.3.1 Effects on the Current Environment**

### **5.3.2 Effects on the Installed Systems**

### **5.3.3 Potential User Problems**

### **5.3.4 Limitations in the Anticipated Implementation Environment That May Inhibit the New Product**

### **5.3.5 Follow-Up Problems**

## **5.4 Tasks**

### **5.4.1 Project Planning**

### **5.4.2 Planning of the Development Phases**

## **5.5 Migration to the New Product**

### **5.5.1 Requirements for Migration of the New Product**

### **5.5.2 Data That Has to Be Modified or Translated for the New System**

## **5.6 Risks**

## **5.7 Costs**

## **5.8 User Documentation and Training**

### **5.8.1 User Documentation Requirements**

### **5.8.2 Training Requirements**

## **5.9 Waiting Room**

## **5.10 Ideas for Solutions**

# Bibliography

- [1] Charades dataset. Allen Institute for Artificial Intelligence. [Online]. Available: <http://allenai.org/plato/charades/>
- [2] J. Xu, T. Mei, T. Yao, and Y. Rui, “Msr-vtt: A large video description dataset for bridging video and language.” IEEE International Conference on Computer Vision and Pattern Recognition (CVPR), June 2016. [Online]. Available: <https://www.microsoft.com/en-us/research/publication/msr-vtt-a-large-video-description-dataset-for-bridging-video-and-language/>
- [3] R. Kiros, R. Salakhutdinov, and R. S. Zemel, “Unifying visual-semantic embeddings with multimodal neural language models,” *CoRR*, vol. abs/1411.2539, 2014. [Online]. Available: <http://arxiv.org/abs/1411.2539>
- [4] T. Pfister, J. Charles, and A. Zisserman, “Flowing convnets for human pose estimation in videos,” *CoRR*, vol. abs/1506.02897, 2015. [Online]. Available: <http://arxiv.org/abs/1506.02897>
- [5] V. Belagiannis and A. Zisserman, “Recurrent human pose estimation,” *CoRR*, vol. abs/1605.02914, 2016. [Online]. Available: <http://arxiv.org/abs/1605.02914>
- [6] S. Wei, V. Ramakrishna, T. Kanade, and Y. Sheikh, “Convolutional pose machines,” *CoRR*, vol. abs/1602.00134, 2016. [Online]. Available: <http://arxiv.org/abs/1602.00134>
- [7] X. Zhou, M. Zhu, K. Derpanis, and K. Daniilidis, “Sparseness meets deepness: 3d human pose estimation from monocular video,” in *CVPR*, 2016.
- [8] D. Holden, J. Saito, and T. Komura, “A deep learning framework for character motion synthesis and editing,” *ACM Trans. Graph.*, vol. 35, no. 4, pp. 138:1–138:11, Jul. 2016. [Online]. Available: <http://doi.acm.org/10.1145/2897824.2925975>